

[Upgrade](#)

Towards
Data Science

[DATA SCIENCE](#)[MACHINE LEARNING](#)[PROGRAMMING](#)[VISUALIZATION](#)[AI](#)[JOURNALISM](#)[EVENTS](#)[SUBMIT](#)

Towards Data Science

Sharing concepts,
ideas, and codes.

[Following](#)

Tara Boyle

[Follow](#)

Feb 4 · 5 min read

Dealing with Imbalanced Data

A guide to effectively handling imbalanced datasets in Python



[Upgrade](#)

Towards
Data Science

[DATA SCIENCE](#)[MACHINE LEARNING](#)[PROGRAMMING](#)[VISUALIZATION](#)[AI](#)[JOURNALISM](#)[EVENTS](#)[SUBMIT](#)

Science

Sharing concepts,
ideas, and codes.

[Following](#)

327



Photo by Ales Nesetril on Unsplash

Imbalanced classes are a common problem in machine learning classification where there are a disproportionate ratio of observations in each class. Class imbalance can be found in many different areas including medical diagnosis, spam filtering, and fraud detection. In this guide, we'll look at five possible ways to handle an imbalanced class problem. Important Note: This guide will focus solely on addressing imbalanced classes and will not addressing other important machine learning steps including, but not limited to, feature selection or hyperparameter tuning.

Data

We will use the Credit Card Fraud Detection Dataset available on Kaggle. The dataset is high imbalanced, with only 0.17% of transactions being classified as fraudulent. The full notebook can be found [here](#).

Our objective will be to correctly classify the minority class of fraudulent transactions.

Distribution of Transactions



The Problem with Imbalanced Classes

Most machine learning algorithms work best when the number of samples in each class are about equal. This is because most algorithms are designed to maximize accuracy and reduce error.

The Problem with Accuracy



Upgrade

Towards
Data Science

DATA SCIENCE

MACHINE LEARNING

PROGRAMMING

VISUALIZATION

AI

JOURNALISM

EVENTS

SUBMIT

Towards Data Science

Sharing concepts, ideas, and codes.

Following



327



```
4
5 # setting up testing and training sets
6 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=27)
7
8 # DummyClassifier to predict only target 0
9 dummy = DummyClassifier(strategy='most_frequent').fit(X_train, y_train)
10 dummy_pred = dummy.predict(X_test)
11
12 # checking unique labels
13 print('Unique predicted labels: ', (np.unique(dummy_pred)))
14
15 # checking accuracy
16 print('Test score: ', accuracy_score(y_test, dummy_pred))
17
18 Unique predicted labels: [0]
19 Test score: 0.9981461194910255
```

dummy_classifier.py hosted with ❤ by GitHub

[view raw](#)

We got an accuracy score of 99.8% — And without even training a model! Let's compare this to logistic regression, an actual trained classifier.

```
1 # Modeling the data as is
2 # Train model
3 lr = LogisticRegression(solver='liblinear').fit(X_train, y_train)
4
5 # Predict on training set
6 lr_pred = lr.predict(X_test)
7
8 # Checking accuracy
9 accuracy_score(y_test, lr_pred)
10 0.9992
11
12 # Checking unique values
13 predictions = pd.DataFrame(lr_pred)
14 predictions[0].value_counts()
15 0    71108
16 1      94
```

lr_baseline.py hosted with ❤ by GitHub

[view raw](#)



Upgrade

Towards
Data Science

DATA SCIENCE

MACHINE LEARNING

PROGRAMMING

VISUALIZATION

AI

JOURNALISM

EVENTS

SUBMIT

Towards Data Science

Sharing concepts,
ideas, and codes.

Following



327



can be very misleading. Metrics that can provide better insight include:

- **Confusion Matrix:** a table showing correct predictions and types of incorrect predictions.
- **Precision:** the number of true positives divided by all positive predictions. Precision is also called Positive Predictive Value. It is a measure of a classifier's exactness. Low precision indicates a high number of false positives.
- **Recall:** the number of true positives divided by the number of positive values in the test data. Recall is also called Sensitivity or the True Positive Rate. It is a measure of a classifier's completeness. Low recall indicates a high number of false negatives.
- **F1: Score:** the weighted average of precision and recall.

Let's see what happens when we apply these F1 and recall scores to our logistic regression from above.

```
1 # f1 score
2 f1_score(y_test, lr_pred)
3 0.7522
4
5 # recall score
6 recall_score(y_test, lr_pred)
7 0.6439
```

change_metric.py hosted with ❤ by GitHub

[view raw](#)

These scores don't look quite so impressive. Let's see what other methods we might try to improve our new metrics.

2. Change the algorithm

While in every machine learning problem, it's a good rule of thumb to try a variety of algorithms, it can be especially beneficial with imbalanced datasets. Decision trees frequently perform well on imbalanced data. They work by learning a hierarchy of if/else questions and this can force both classes to be addressed.

```
1 from sklearn.ensemble import RandomForestClassifier
2
3 # train model
4 rfc = RandomForestClassifier(n_estimators=10).fit(X_train, y_train)
5
6 # predict on test set
7 rfc_pred = rfc.predict(X_test)
8
9 accuracy_score(y_test, rfc_pred)
10 0.9995
11
```



Upgrade

Towards
Data Science

DATA SCIENCE

MACHINE LEARNING

PROGRAMMING

VISUALIZATION

AI

JOURNALISM

EVENTS

SUBMIT

**Towards Data
Science**Sharing concepts,
ideas, and codes.

Following



327



change_algorithm.py hosted with ❤ by GitHub

[view raw](#)

While our accuracy score is slightly lower, both F1 and recall have increased as compared to logistic regression! It appears that for this specific problem, random forest may be a better choice of model.

3. Resampling Techniques — Oversample minority class

Our next method begins our resampling techniques.

Oversampling can be defined as adding more copies of the minority class. Oversampling can be a good choice when you don't have a ton of data to work with.

We will use the resampling module from Scikit-Learn to randomly replicate samples from the minority class.

Important Note

Always split into test and train sets BEFORE trying oversampling techniques! Oversampling before splitting the data can allow the exact same observations to be present in both the test and train sets.

This can allow our model to simply memorize specific data points and cause overfitting and poor generalization to the test data.

```
1  from sklearn.utils import resample
2
3  # Separate input features and target
4  y = df.Class
5  X = df.drop('Class', axis=1)
6
7  # setting up testing and training sets
8  X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=27)
9
10 # concatenate our training data back together
11 X = pd.concat([X_train, y_train], axis=1)
12
13 # separate minority and majority classes
14 not_fraud = X[X.Class==0]
15 fraud = X[X.Class==1]
16
17 # upsample minority
18 fraud_upsampled = resample(fraud,
19                             replace=True, # sample with replacement
20                             n_samples=len(not_fraud), # match number in majority class
21                             random_state=27) # reproducible results
22
23 # combine majority and upsampled minority
```



Upgrade

Towards
Data Science

DATA SCIENCE

MACHINE LEARNING

PROGRAMMING

VISUALIZATION

AI

JOURNALISM

EVENTS

SUBMIT

Towards Data Science

Sharing concepts,
ideas, and codes.

Following



327



upsampled.py hosted with ❤ by GitHub

[view raw](#)

After resampling we have an equal ratio of data points for each class! Let's try our logistic regression again with the balanced training data.

```
1 # trying logistic regression again with the balanced dataset
2 y_train = upsampled.Class
3 X_train = upsampled.drop('Class', axis=1)
4
5 upsampled = LogisticRegression(solver='liblinear').fit(X_train, y_train)
6
7 upsampled_pred = upsampled.predict(X_test)
8
9 # Checking accuracy
10 accuracy_score(y_test, upsampled_pred)
11 0.9807
12
13 # f1 score
14 f1_score(y_test, upsampled_pred)
15 0.1437
16
17 recall_score(y_test, upsampled_pred)
18 0.8712
```

upsampled_lr.py hosted with ❤ by GitHub

[view raw](#)

Our recall score increased, but F1 is much lower than with either our baseline logistic regression or random forest from above. Let's see if undersampling might perform better here.

4. Resampling techniques — Undersample majority class

Undersampling can be defined as removing some observations of the majority class. Undersampling can be a good choice when you have a ton of data -think millions of rows. But a drawback is that we are removing information that may be valuable. This could lead to underfitting and poor generalization to the test set.

We will again use the resampling module from Scikit-Learn to randomly remove samples from the majority class.

```
1 # still using our separated classes fraud and not_fraud from above
2
3 # downsample majority
4 not_fraud_downsampled = resample(not_fraud,
5                                  replace = False, # sample without replacement
```



Upgrade

Towards
Data Science

DATA SCIENCE

MACHINE LEARNING

PROGRAMMING

VISUALIZATION

AI

JOURNALISM

EVENTS

SUBMIT

Towards Data Science

Sharing concepts, ideas, and codes.

Following



327



```

11
12 # checking counts
13 downsamped.Class.value_counts()
14      1      360
15      0      360

```

downsamped.py hosted with ❤ by GitHub

[view raw](#)

Again, we have an equal ratio of fraud to not fraud data points, but in this case a much smaller quantity of data to train the model on. Let's again apply our logistic regression.

```

1 # trying logistic regression again with the undersampled dataset
2
3 y_train = downsamped.Class
4 X_train = downsamped.drop('Class', axis=1)
5
6 undersampled = LogisticRegression(solver='liblinear').fit(X_train, y_train)
7
8 undersampled_pred = undersampled.predict(X_test)
9
10 # Checking accuracy
11 accuracy_score(y_test, undersampled_pred)
12      0.9758
13
14 # f1 score
15 f1_score(y_test, undersampled_pred)
16      0.1171
17
18 recall_score(y_test, undersampled_pred)
19      0.8636

```

downsamped_lr.py hosted with ❤ by GitHub

[view raw](#)

Undersampling underperformed oversampling in this case. Let's try one more method for handling imbalanced data.

5. Generate synthetic samples

A technique similar to upsampling is to create synthetic samples. Here we will use imblearn's SMOTE or Synthetic Minority Oversampling Technique. SMOTE uses a nearest neighbors algorithm to generate new and synthetic data we can use for training our model. Again, it's important to generate the new samples only in the training set to ensure our model generalizes well to unseen data.

```

1 from imblearn.over sampling import SMOTE

```



Upgrade

Towards
Data Science

DATA SCIENCE

MACHINE LEARNING

PROGRAMMING

VISUALIZATION

AI

JOURNALISM

EVENTS

SUBMIT

Towards Data Science

Sharing concepts,
ideas, and codes.

Following



327



```
7 # setting up testing and training sets
8 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=27)
9
10 sm = SMOTE(random_state=27, ratio=1.0)
11 X_train, y_train = sm.fit_sample(X_train, y_train)
```

smote.py hosted with ❤ by GitHub

[view raw](#)

After generating our synthetic data points, let's see how our logistic regression performs.

```
1 smote = LogisticRegression(solver='liblinear').fit(X_train, y_train)
2
3 smote_pred = smote.predict(X_test)
4
5 # Checking accuracy
6 accuracy_score(y_test, smote_pred)
7     0.9858
8
9 # f1 score
10 f1_score(y_test, smote_pred)
11     0.1846
12
13 recall_score(y_test, smote_pred)
14     0.8636
```

smote_lr.py hosted with ❤ by GitHub

[view raw](#)

Our F1 score is increased and recall is similar to the upsampled model above and for our data here outperforms undersampling.

Conclusion

We explored 5 different methods for dealing with imbalanced datasets:

1. Change the performance metric
2. Change the algorithm
3. Oversample minority class
4. Undersample majority class
5. Generate synthetic samples

It appears for this particular dataset random forest and SMOTE are among the best of the options we tried here.

These are just some of the many possible methods to try when dealing with imbalanced datasets, and not an exhaustive list. Some other methods to consider are collecting more data or choosing different resampling ratios — you don't have to have exactly a 1:1 ratio! You should always try several approaches and then decide which is best for your problem.

[Upgrade](#)

Towards
Data Science

[DATA SCIENCE](#)[MACHINE LEARNING](#)[PROGRAMMING](#)[VISUALIZATION](#)[AI](#)[JOURNALISM](#)[EVENTS](#)[SUBMIT](#)

Towards Data Science

Sharing concepts, ideas, and codes.

[Following](#)

327



WRITTEN BY

Tara Boyle

[Follow](#)

I'm passionate about all things data! I'm interested in applying machine learning to create business solutions.



Towards Data Science

Sharing concepts, ideas, and codes.

[Following](#)[See responses \(3\)](#)

More From Medium

[More from Towards Data Science](#)[More from Towards Data Science](#)[More from Towards Data Science](#)



Upgrade



Towards
Data Science

DATA SCIENCE

MACHINE LEARNING

PROGRAMMING

VISUALIZATION




AI

JOURNALISM

EVENTS

SUBMIT

Learning Engineer



 **John D. Bourke** in Towards Data Science
Jul 6 · 11 min read ★  5.6K | 

Sharing concepts,
ideas, and codes.

Following

 327



Susan Li in Towards Data Science
Jul 17 · 10 min read  896 | 



Admond Lee in Towards Data Science
Jul 13 · 11 min read ★  908 | 